

Homework 1

Elliott Pryor

30 September 2023

Problem 1**Statement:**

Consider the problem of imitation learning within a discrete MDP with horizon T and an expert policy π^* . We gather expert demonstrations from π^* and fit an imitation policy π_θ to these trajectories so that

$$\mathbb{E}_{p_{\pi^*}(s)} \pi_\theta(a \neq \pi^*(s)) = \frac{1}{T} \sum_{t=1}^T \mathbb{E}_{p_{\pi^*}(s_t)} \pi_\theta(a \neq \pi^*(s_t)) \leq \epsilon \quad (1)$$

i.e., the expected likelihood that the learned policy π_θ disagrees with the expert π^* within the training distribution p_{π^*} of states drawn from random expert trajectories is at most ϵ . For convenience, the notation $p_\pi(s_t)$ indicates the state distribution under π at time step t while $p(s)$ indicates the state marginal of π across time steps, unless indicated otherwise

1. Show that $\sum_{s_t} |p_{\pi_\theta}(s_t) - p_{\pi^*}(s_t)| \leq 2T\epsilon$
2. Consider the expected return of the learned policy π_θ for a state-dependent reward function $r(s_t)$. where we assume the reward is bounded $|r(s_t)| \leq R_{max}$:

$$J(\pi) = \sum_{t=1}^T \mathbb{E}_{p_\pi(s_t)} r(s_t)$$

- (a) Show that $|J(\pi_\theta) - J(\pi^*)| \in \mathcal{O}(T\epsilon)$ when the reward only depends on the last state. I.e., $r(s_t) = 0 \ \forall t < T$
- (b) Show that $J(\pi^*) - J(\pi_\theta) \in \mathcal{O}(T^2\epsilon)$

Solution

1. So we want to show that $\sum_{s_t} |p_{\pi_\theta}(s_t) - p_{\pi^*}(s_t)| \leq 2T\epsilon$. We know that $\mathbb{E}_{p_{\pi^*}(s)} \pi_\theta(a \neq \pi^*(s)) = \frac{1}{T} \sum_{t=1}^T \mathbb{E}_{p_{\pi^*}(s_t)} \pi_\theta(a \neq \pi^*(s_t)) \leq \epsilon$. This means that the expected value of the probability that π_θ disagrees with π^* is at most ϵ .

We note that if there is no mistake, then there is no difference in the state distribution. In the worst case, after a mistake the state distribution is completely different, so $|p_{\pi_\theta}(s_t) - p_{\pi^*}(s_t)| = 2 \ \forall t > t_{mistake}$. Thus the cost after the mistake would be: $2(T - t_{mistake})$. We partition the space by the time of the mistake: $\epsilon * cost_1, (1 - \epsilon)\epsilon * cost_2, (1 - \epsilon)^2\epsilon * cost_3, \dots, (1 - \epsilon)^T * 0$. Where cost is the total difference in the state distribution after the mistake. So all the possible costs that we can incur are: $\epsilon * 2T, (1 - \epsilon)\epsilon * 2(T - 1), (1 - \epsilon)^2\epsilon * 2(T - 2), \dots, (1 - \epsilon)^T * 0$. So the worst case is we make a mistake on the very first sample. Thus $\sum_{s_t} |p_{\pi_\theta}(s_t) - p_{\pi^*}(s_t)| \leq 2T\epsilon$

2. (a) We want to show that $|J(\pi_\theta) - J(\pi^*)| \in \mathcal{O}(T\epsilon)$ when the reward only depends on the last state. This bound is somewhat misleading as it appears to grow linearly, but in fact is constant. This is because the reward only depends on the last state, so $|J(\pi_\theta) - J(\pi^*)| = |\mathbb{E}_{p_{\pi^*}(s_T)}(r(s_T)) - \mathbb{E}_{p_{\pi_\theta}(s_T)}(r(s_T))|$. We know that $r(s)$ is a bounded function, so this is clearly $\leq R_{max}$ which is constant.

We can also derive the original (linear) bound fairly easily. We partition it into the case of no mistakes and the case of mistakes. No mistake occurs with probability $(1 - \epsilon)^T$ which by Bernoulli's inequality is $\geq 1 - T\epsilon$. This means the probability of a mistake is $\leq 1 - (1 - T\epsilon) = T\epsilon$. We can assume worst case cost if a mistake occurs. Thus the expected value of the cost $\leq (1 - T\epsilon) * 0 + (T\epsilon) * R_{max} \in \mathcal{T}\epsilon$

- (b) Given the previous result, we consider the generic case where the reward depends on all states, as a sum of cases where the reward only depends on the last state. Let $j(T)$ denote the cost discovered in the previous section $j(t) = t\epsilon * R_{max}$. Then the generic case:

$$\begin{aligned}
 |J(\pi^*) - J(\pi_\theta)| &\leq j(1) + j(2) + \dots + j(T) \\
 &= \epsilon * R_{max} * (1 + 2 + \dots + T) \\
 &= \epsilon * R_{max} * \frac{T(T+1)}{2} \\
 &= \epsilon * R_{max} * \left(\frac{T^2}{2} + \frac{T}{2} \right) \\
 &\in \mathcal{O}(T^2\epsilon)
 \end{aligned}$$

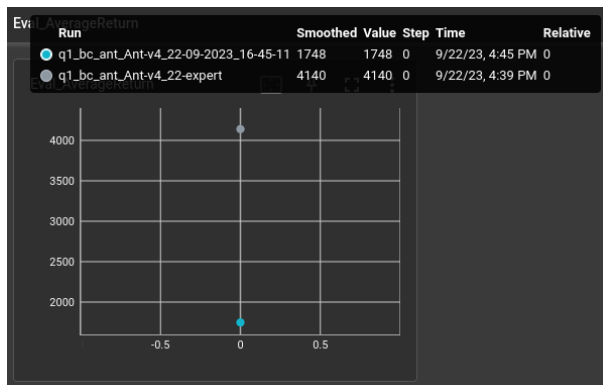
Problem 2

Statement:

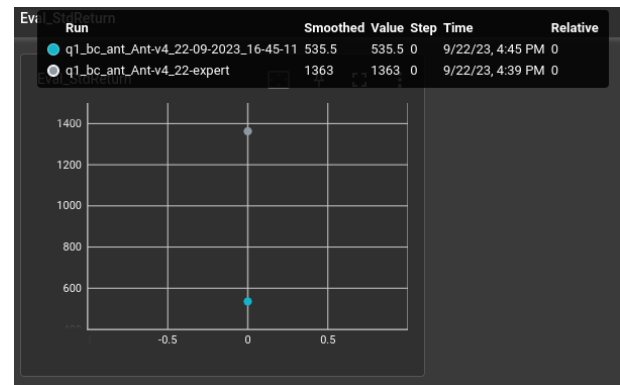
Run behavioral cloning (BC) and report results on two tasks: the Ant environment, where a behavioral cloning agent should achieve at least 30% of the performance of the expert, and one environment of your choosing where it does not.

Solution

Figure 1 show the results on the ant environment with default hyperparameters. 2 layers, of 64 nodes, 0.005 lr, batch size = 1000, train steps = 1000. Metrics were computed over 5 rollouts (Eval batch size = 5000, ep len = 1000).



(a) Mean Returns



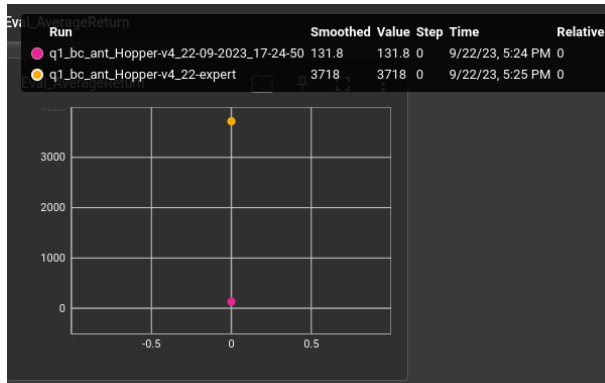
(b) Std of Returns

Figure 1: Ant environment with 5 repeats (eval batch size = 5000, batch size = 1000). Blue is expert, grey is BC.

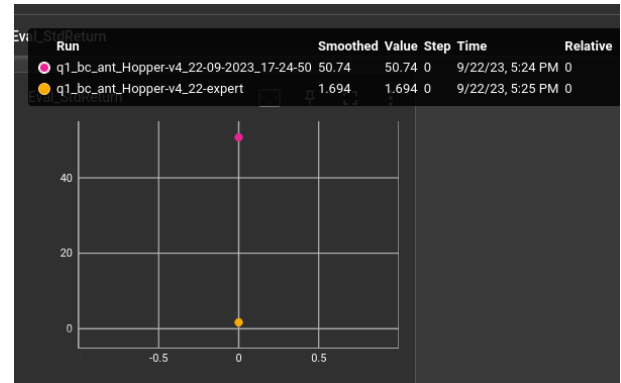
Figure 2 show the results on the hopper environment with default hyperparameters. We see much worse performance here, indicating that the BC agent did not work as well. This could be because the environment is much more challenging, and possibly a larger network would be needed.

We try double the network depth in case the network is not large enough to capture the function, and 5 times the batch size in order to give it more training data (we also proportionally increase the training iterations), and see the results in Figure 3. They did not improve the performance, and in fact made it worse.

Based on this, we vary batch size more. The results can be seen in figure 4, the effect is not very significant, and we do not recover 30% of the expert performance.



(a) Mean Returns



(b) Std of Returns

Figure 2: Hopper environment with 5 repeats (eval batch size = 5000, batch size = 1000). Yellow is expert, Pink is BC.

Run	Value	Step	Time	Relative
q1_bc_ant_Hopper-v4_22-batch5K	247.8	0	9/25/23, 9:41 AM	0
q1_bc_ant_Hopper-v4_22-default	131.8	0	9/22/23, 5:24 PM	0
q1_bc_ant_Hopper-v4_22-expert	3718	0	9/22/23, 5:25 PM	0

Figure 3: Different networks on hopper. This is the mean returns

Run	Value	Step	Time	Relative
q1_bc_ant_Hopper-v4_22-batch10K	266.1	0	9/25/23, 10:16 AM	0
q1_bc_ant_Hopper-v4_22-batch15K	273.9	0	9/25/23, 10:17 AM	0
q1_bc_ant_Hopper-v4_22-batch20K	265.5	0	9/25/23, 10:18 AM	0
q1_bc_ant_Hopper-v4_22-batch5K	247.8	0	9/25/23, 9:41 AM	0
q1_bc_ant_Hopper-v4_22-default	131.8	0	9/22/23, 5:24 PM	0
q1_bc_ant_Hopper-v4_22-expert	3718	0	9/22/23, 5:25 PM	0

Figure 4: Variation of batch size and training iterations (batch = train) on hopper. This is the mean returns.

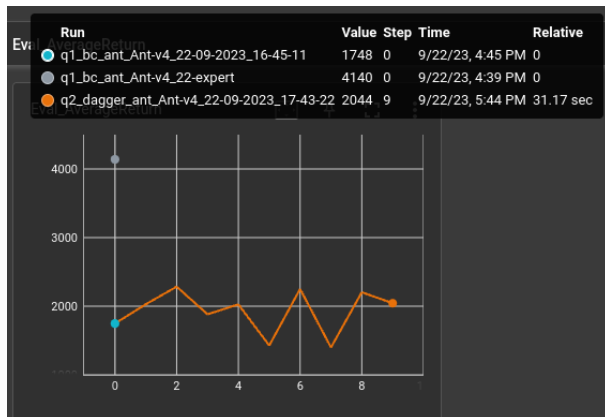
Problem 3

Statement:

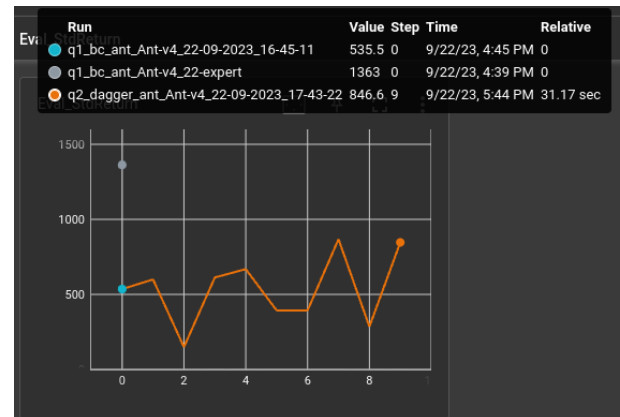
Run DAgger and report results on the two tasks you tested previously with behavioral cloning (i.e., Ant + another environment).

Solution

Figure 5 show the results on the ant environment with default hyperparameters. 2 layers, of 64 nodes, 0.005 lr. Batch size = 1000, and 10 dagger iterations. Metrics were computed over 5 rollouts.



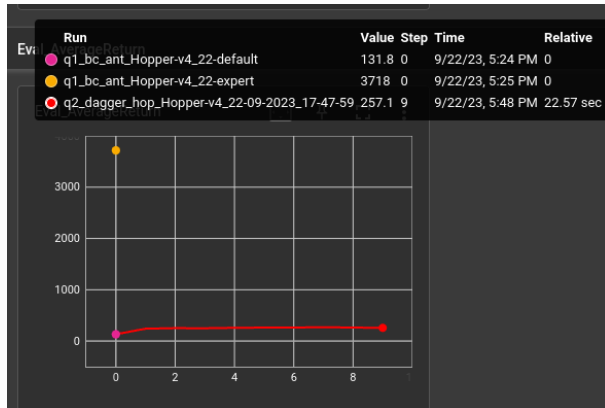
(a) Mean Returns



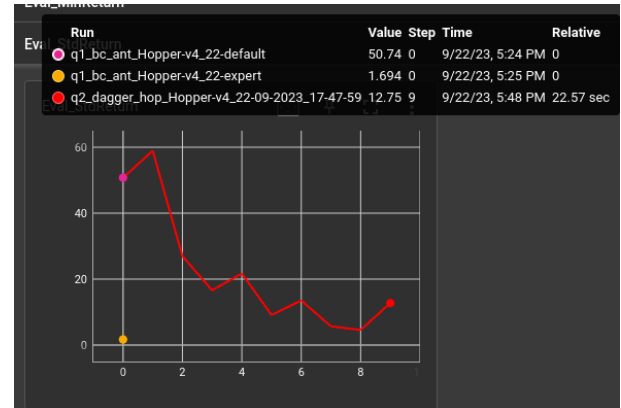
(b) Std of Returns

Figure 5: Ant environment with 5 repeats (eval batch size = 5000, batch size = 1000). Blue is expert, grey is BC, orange is DAGger.

We also tried doubling the network depth. And the average returns can be seen in Figure 7. Once again, the larger network does not seem to do any better.



(a) Mean Returns



(b) Std of Returns

Figure 6: hop environment with 5 repeats (eval batch size = 5000, batch size = 1000). Yellow is expert, Pink is BC, Red is DAGger.

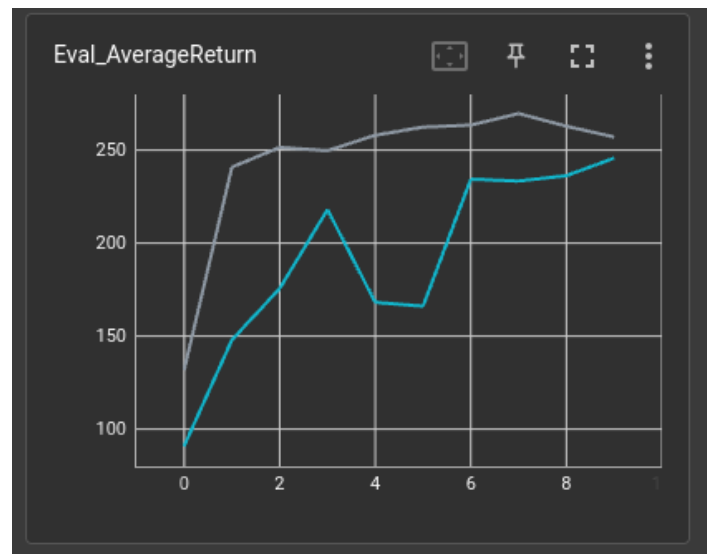


Figure 7: Comparison of dagger with default network size (grey) and double network depth (blue) in average return.